



Training versus validation loss — faint: unsmoothed; bold: 10-epoch mean


01 · Baseline (patience 30)




02 · Constant LR (patience 100)




02 · Plateau LR (patience 100)



Later training (epochs 26 onward)



Later training (epochs 26 onward)



Later training (epochs 26 onward)

